

This tutorial is modified from: <https://www.dataquest.io/blog/jupyter-notebook-tutorial/> (<https://www.dataquest.io/blog/jupyter-notebook-tutorial/>)

In [1]:

```
print('Hello World!')
```

Hello World!

In [2]:

```
import time
time.sleep(5)
```

In [3]:

```
def say_hello(recipient):
    return 'Hello, {}'.format(recipient)
say_hello('EE4211 Students')
```

Out[3]:

'Hello, EE4211 Students!'

Intro to Markdown (This is a level 1 heading)

This is a level 2 heading

This is some plain text that forms a paragraph. Add emphasis via **bold** and **bold**, or italic and italic. Paragraphs must be separated by an empty line.

- Sometimes we want to include lists.
 - Which can be indented.
- 1. Lists can also be numbered.
- 2. It is possible to include hyperlinks (<https://www.ece.nus.edu.sg/stfpage/motani/>)
- 3. You can add images too:



In [4]:

```
import numpy as np
def square(x):
    return x * x
```

In [5]:

```
x = np.random.randint(1, 10)
y = square(x)
print('%d squared is %d' % (x, y))
```

4 squared is 16

Let get some data from Fortune 500

(http://archive.fortune.com/magazines/fortune/fortune500_archive/full/2005/) companies spanning over 50 years since the list's first publication in 1955, put together from Fortune's public archive. You can get a CSV of the data at this link (<https://s3.amazonaws.com/dq-blog-files/fortune500.csv>).

In [6]:

```
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
sns.set(style="darkgrid")
```

In [7]:

```
df = pd.read_csv('fortune500.csv')
```

In [8]:

```
df.head()
```

Out[8]:

	Year	Rank	Company	Revenue (in millions)	Profit (in millions)
0	1955	1	General Motors	9823.5	806
1	1955	2	Exxon Mobil	5661.4	584.8
2	1955	3	U.S. Steel	3250.4	195.4
3	1955	4	General Electric	2959.1	212.6
4	1955	5	Esmark	2510.8	19.1

In [9]:

```
df.tail()
```

Out[9]:

	Year	Rank	Company	Revenue (in millions)	Profit (in millions)
25495	2005	496	Wm. Wrigley Jr.	3648.6	493
25496	2005	497	Peabody Energy	3631.6	175.4
25497	2005	498	Wendy's International	3630.4	57.8
25498	2005	499	Kindred Healthcare	3616.6	70.6
25499	2005	500	Cincinnati Financial	3614.0	584

In [10]:

```
#Let's rename those columns so we can refer to them later.  
df.columns = ['year', 'rank', 'company', 'revenue', 'profit']
```

In [11]:

```
len(df)
```

Out[11]:

25500

In [12]:

```
df.dtypes
```

Out[12]:

```
year          int64  
rank          int64  
company       object  
revenue       float64  
profit        object  
dtype: object
```

In [13]:

```
non_numeric_profits = df.profit.str.contains('[^0-9.-]')  
df.loc[non_numeric_profits].head()
```

Out[13]:

	year	rank	company	revenue	profit
228	1955	229	Norton	135.0	N.A.
290	1955	291	Schlitz Brewing	100.0	N.A.
294	1955	295	Pacific Vegetable Oil	97.9	N.A.
296	1955	297	Liebmann Breweries	96.0	N.A.
352	1955	353	Minneapolis-Moline	77.4	N.A.

In [14]:


```
# How much bad data do we have?  
round(100*len(df.profit[non_numeric_profits])/len(df),2)
```

Out[14]:

1.45

In [15]:

```
bin_sizes, _, _ = plt.hist(df.year[non_numeric_profits], bins=range(1955, 2006))
```



In [16]:

```
#Let's remove the bad rows
df = df.loc[~non_numeric_profits]
df.profit = df.profit.apply(pd.to_numeric)
```

In [17]:

```
len(df)
```

Out[17]:

25131

In [18]:

```
df.dtypes
```

Out[18]:

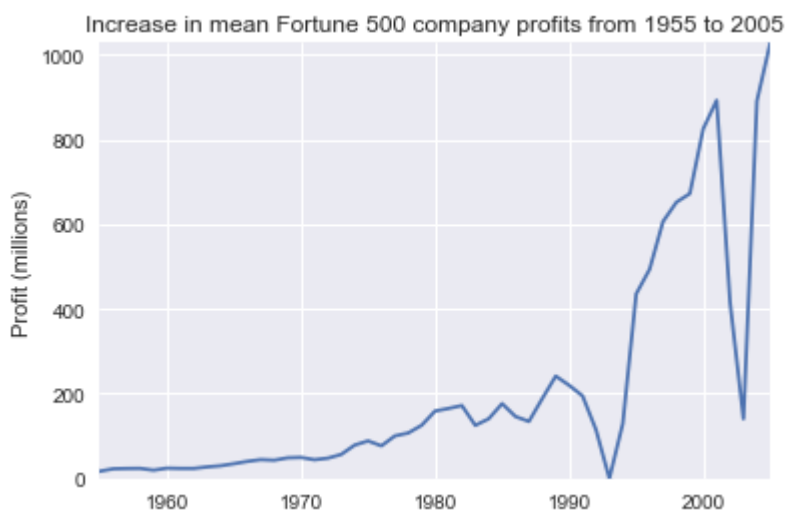
```
year          int64
rank          int64
company       object
revenue       float64
profit        float64
dtype: object
```

In [19]:

```
group_by_year = df.loc[:, ['year', 'revenue', 'profit']].groupby('year')
avgs = group_by_year.mean()
x = avgs.index
def plot(x, y, ax, title, y_label):
    ax.set_title(title)
    ax.set_ylabel(y_label)
    ax.plot(x, y)
    ax.margins(x=0, y=0)
```

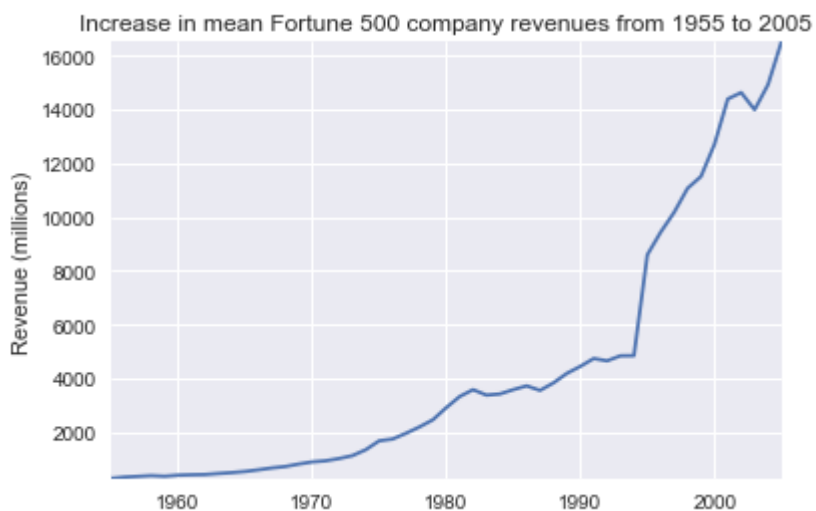
In [20]:

```
# Let's plot profits
y1 = avgs.profit
fig, ax = plt.subplots()
plot(x, y1, ax, 'Increase in mean Fortune 500 company profits from 1955 to 2005',
    'Profit (millions)')
```



In [21]:

```
# Let's plot revenues
y2 = avgs.revenue
fig, ax = plt.subplots()
plot(x, y2, ax, 'Increase in mean Fortune 500 company revenues from 1955 to 2005',
    'Revenue (millions)')
```



In [22]:

```
# A more complex plot, code taken from StackOverflow: https://stackoverflow.com/a/47582329/604687
def plot_with_std(x, y, stds, ax, title, y_label):
    ax.fill_between(x, y - stds, y + stds, alpha=0.2)
    plot(x, y, ax, title, y_label)

fig, (ax1, ax2) = plt.subplots(ncols=2)
title = 'Increase in mean and std Fortune 500 company %s from 1955 to 2005'
stds1 = group_by_year.std().profit.values
stds2 = group_by_year.std().revenue.values
plot_with_std(x, y1.values, stds1, ax1, title % 'profits', 'Profit (millions)')
plot_with_std(x, y2.values, stds2, ax2, title % 'revenues', 'Revenue (millions)')
fig.set_size_inches(14, 4)
fig.tight_layout()
```

